

# TDDC17 Artificial Intelligence

## Lab 1: Intelligent Agents

### Aim

The "agent" paradigm has rapidly become one of the major conceptual frameworks for understanding and structuring research activities within Artificial Intelligence. In addition, many software frameworks are being introduced and used both in research and industry to control the complexity of sophisticated distributed software systems where the nodes exhibit behaviors often associated with intelligent behavior.

The purpose with this lab is to introduce the agent concept by allowing you to familiarize yourself with a software agent environment and world simulator. You will program an agent and implement its behaviour to provide cleaning capabilities in the vacuum cleaning world. The behaviour will then be compared to some simple reference agents.

### Preparation

Read chapter 2 in the course book. Pay particular attention to section 2.3 (agent environments) and section 2.4 (different agent types). After this, read this lab description in detail.

### Lab assignment

Your assignment is to construct a vacuum agent. The agent should be able to remove all the dirt in a rectangular world of varying size with a random distribution of dirt. When the world is cleaned the agent should return to the start position and shut down.

### Lab system

Starting 2019 two alternative implementations of this lab are provided: Java and Python. Note this pertains only to Lab 1. The Java implementation of this lab has been used and refined for many years and is still the default version. In contrast, the Python implementation is completely new and less extensively tested. If you choose to use Python, please report any problems you detect to **Mariusz**. The Python implementation is very close to the Java version both in terms of code structure and provided GUI, i.e. you can follow the same instructions below although they mostly refer to the default Java implementation.

The lab system consists of the environment simulator and the agent. You are given a simple example of a vacuum cleaner agent called `MyVacuumAgent.java` **here** (in Python: `myvacuumagent.py` - **here**). What you have to do is to improve on the code controlling the behaviour of the agent to make it perform better. The description of how to set up and run the code is available towards the bottom of the page.

#### The environment simulator

You can test your agent in *the simulator GUI* program. The simulator loads an agent (yours will be implemented in `MyVacuumAgent`) and executes it in a generated environment. Agents execute actions in the environment depending on what they can observe (perceive). The actions influence the environment which in turn generates new sensory impressions for the agents. The simulator handles this interaction in a discrete manner; the world advances in steps, where each step can be described as follows:

1. The agent perceives the world
2. The agent decides on an action
3. The action is executed

The agent should be autonomous and therefore the simulator makes no assumptions about an agent except that it can produce a reaction based on the sensory inputs presented to it. The agents make no assumptions about the simulator's implementation either. This disengagement is important, and makes it possible to compare and evaluate different agents' behaviors. Note that an agent has no authorization to change the world directly, for example via some global reassignment. This disengagement keeps the code modular and independent as it is possible to change the representation of the world but use the same agent.

#### The implementation of the simulator

The simulator supplied for this lab is based on the AIMA-Java/AIMA-Python framework but it is not identical to the original version available on the web.

#### The interface

Basically the agents may perceive dirt or obstacles in the squares, and they can act by moving between the squares, and removing dirt.

#### Percepts

A percept for the vacuum-agent consists of three values:

1. `bump` equal to `true` if the agent has bumped into an obstacle, else `false`.
2. `dirt` equal to `true` if the current square contains dirt, else `false`.
3. `home` equal to `true` if the agent is at the start position, else `false`.

In Java the three Boolean values are obtained with the following code located in the `execute` method of an agent:

```
DynamicPercept p = (DynamicPercept) percept;  
Boolean bump = (Boolean)p.getAttribute("bump");  
Boolean dirt = (Boolean)p.getAttribute("dirt");  
Boolean home = (Boolean)p.getAttribute("home");
```

Similarly in Python implementation those values are obtained with the following code:

```
bump = percept.attributes["bump"]  
dirt = percept.attributes["dirt"]  
home = percept.attributes["home"]
```

Note that this is the only data that the agent gets from the simulator.

#### Output

In Java the output from the agent's program (return value) is one of the following actions:

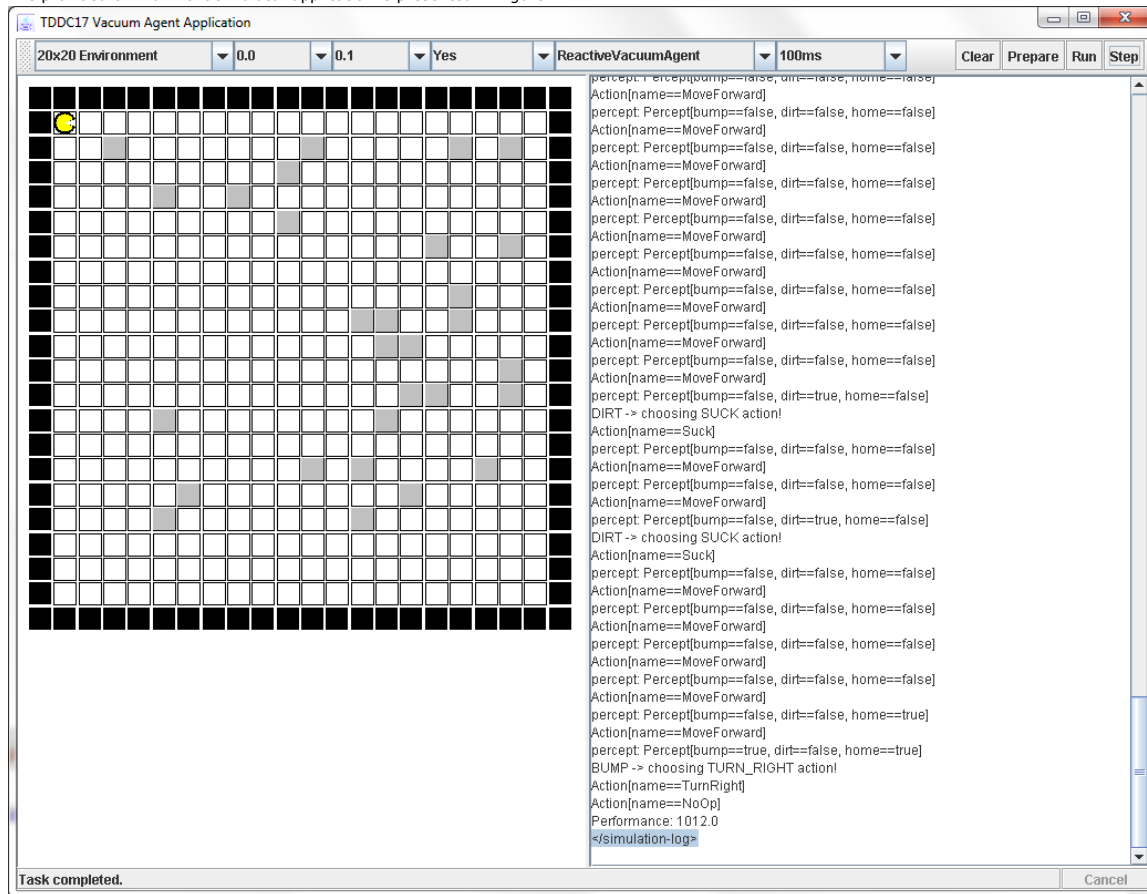
1. `LIUVacuumEnvironment.ACTION_MOVE_FORWARD` moves the agent one step forward depending on its direction.

2. `LIUVacuumEnvironment.ACTION_TURN_LEFT` makes the agent turn left.
3. `LIUVacuumEnvironment.ACTION_TURN_RIGHT` makes the agent turn right.
4. `LIUVacuumEnvironment.ACTION_SUCK` cleans the dirt from the agent's current location.
5. `NoOpAction.NO_OP` informs the simulator that the agent has finished cleaning.

The Python version uses one of the following actions: `ACTION_FORWARD`, `ACTION_TURN_LEFT`, `ACTION_TURN_RIGHT`, `ACTION_SUCK`, `ACTION_NOP`, respectively.

## GUI

The provided environment simulator application is presented in Figure 1.



**Figure 1:** Example of an environment in the vacuum agents' world.

The main window of the application is divided into two parts: the environment visualization pane on the left and the text output pane which displays the redirected `System.out` in Java and `self.log` in Python. Certain messages from the simulator are printed in this pane but it is also useful to debug your agent's code.

In the environment pane the black squares represent obstacles. Trying to move an agent into an obstacle results in perceiving the `bump` equal to `true`. The grey squares represent the dirt. Moving into such a square results in perceiving the `dirt` equal to `true`. The graphical representation of the agent visualizes its current direction.

The following drop down menus are available to set the simulation parameters (from left):

1. Environment size can be set to 5x5, 10x10, 15x15, 20x20, 5x10, and 10x5. Your agent should not be tailored to any specific size, though for convenience you can assume a maximum size of 20x20.
2. Hinder probability indirectly specifies the amount of obstacles in the environment (it should be set to 0.0 and 0.1, in Tasks 1 and 2 respectively).
3. Dirt probability indirectly specifies the amount of dirt in the environment.
4. Random environment specifies whether a new environment should be generated for every run. This function is useful to test an agent in the same environment several times.
5. Type of agent lists the available agent type classes. The first on the list is the agent you will modify.
6. Delay specifies the delay between each simulation step.

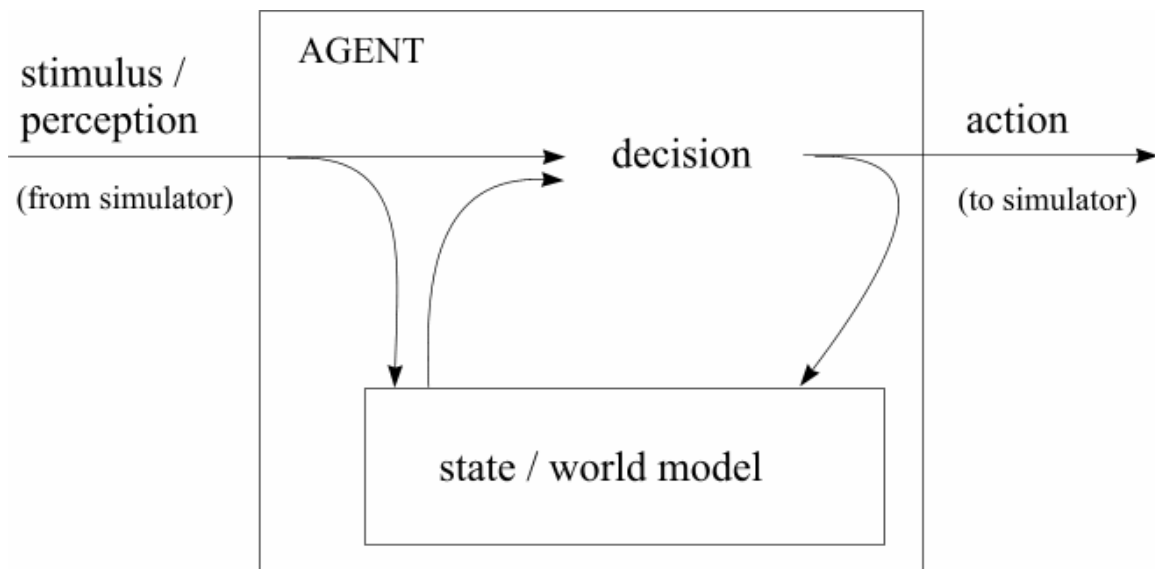
The following buttons are available (from left):

1. Clear cleans the text output pane.
2. Prepare generates a random environment based on the selected parameters.
3. Run starts the simulation.
4. Step allows for performing one step of the simulation at a time. Make sure to execute Prepare first.
5. Cancel terminates the simulation.

## Program

The interesting part of the agent is the `execute` method. It is a function that takes a `percept` as input and returns an action as output. The simulator computes what the agent can perceive (a `percept`), applies the function (your program) to the `percept`, and executes the resulting action. Note (again) that the agent has no authorization to change the world directly, for example via some global reassignment.

As you can see from the definitions in the source code, the `RandomVacuumAgent` and `ReactiveVacuumAgent` (in Java - [here](#), in Python - [here](#)) do not remember what they have done and experienced. If we want the agent to have a memory we have to store information "inside" the agent.



**Figure 2: Agent with internal state**

The state can be a model of the surrounding world as well as some other knowledge that the agent has, e.g. plans.

### Scoring

The performance of an agent is measured by a score. The initial score of the agent is -1000, and for each action the score is changed. Sucking up dirt increases the score 100 points, while shutting down at the starting square gives it 1000 points. All other actions decrease the score by one point.

## Running the Lab environment

You can choose either of the two following lab setups, depending on if you want to use a pre-packaged Eclipse environment or your own favorite editor. If you are unsure, go with Eclipse. Note, the general information about working on this lab remotely can be found [here](#).

### Java Eclipse version (RDP, ThinLinc or on campus)

1. To prepare the local files run the following commands in the terminal:

```

mkdir TDDC17_lab1; cd TDDC17_lab1
cp -r /courses/TDDC17/labs/lab1_workspace/ .

```

2. Setup and start the Eclipse editor:

```

To add the Eclipse editor to your environment execute the following command in the terminal: module add prog/eclipse or alternatively module initadd prog/eclipse to
always add it when logging in - note that you will need to open a new terminal window for the command to take effect the first time. If eclipse is still not available (next step) -
try to force the current available version by using module add prog/eclipse/2023-03 command instead.
eclipse

```

3. When the eclipse editor asks for a path to the workspace choose the lab1\_workspace folder which you copied in step 1.
4. Click *Continue* if you get an *Older Workspace Version* warning.
5. Right-click on the project and select refresh.
6. The Eclipse editor will compile the java files automatically. You can start the lab by using the dropdown menu on the run button (green circle with a white triangle/play symbol). You may have to select "LIUVacuumAgent" in the drop-down menu.

### Java console version (RDP, ThinLinc or on campus)

1. To prepare the local files run the following commands in the terminal:

```

mkdir TDDC17_lab1; cd TDDC17_lab1
mkdir lab1; cd lab1
cp -r /courses/TDDC17/labs/lab1_workspace/project/* .

```

2. To compile the required local classes:

```

./compile

```

3. Now you can run the simulation environment GUI:

```

./start

```

### Python version (console and PyCharm; RDP, ThinLinc or on campus)

1. To prepare the local files run the following command in the terminal:

```

cp -r /courses/TDDC17/labs/TDDC17_lab1_python/ .

```

2. Console: to run the code execute the following command in the TDDC17\_lab1\_python directory:

```

python3 run_lab1.py

```

3. PyCharm: setup and start the PyCharm editor:

```

To add the PyCharm editor to your environment execute the following command in the terminal: module add prog/pycharm-professional/2022.3.1 or alternatively module
initadd prog/pycharm-professional/2022.3.1 to always add it when logging in - note that you will need to open a new terminal window for the command to take effect the
first time.
pycharm.sh
After starting the PyCharm editor, open the project by selecting File/Open and choose the TDDC17_lab1_python directory.
To execute run run_lab1.py - open the file in the editor, right click and choose Run run_lab1.

```

## Running it on your personal computer

You can download the Eclipse compatible workspace here: **.zip file**. Simply extract the file and point Eclipse to it during the start-up ("Select a directory as workspace") and follow the general instructions above. If you do not have Eclipse installed it is recommended that you download a version suitable to your platform **here**.

The Python code is available **here**. You can use your editor of choice (e.g. PyCharm) and follow the instructions above. Note, the lab code requires Python version lower than 3.10 and depends on three additional modules: numpy, tkinter and ipython.

Installation of Python and running the lab on Windows computers:

1. Install a Python distribution which can be downloaded **here**.
2. Open Command Prompt application. Install additional Python packages required by executing following command:

```
pip install numpy ipython
```

In case the path to the pip installation cannot be found and you get the error message you can go to the location of the pip application first before executing the command above:

```
cd C:\Users\YOUR_USER_NAME\AppData\Local\Programs\Python\Python38-32\Scripts
```

3. Start IDLE (Python Shell) application and open run\_lab1.py file which can be found in the downloaded zip file. Press F5 to run the module. Alternatively, you can use other editor of your choice e.g. PyCharm.

In some cases when running the Python version of the lab on MacOS, the graphics may not be displayed properly. If that is the case, you can download an alternative version of the code **here**.

## Detailed Lab Task Description

Play around with the GUI and try the different options and agents.

1. Using `MyVacuumAgent.java` (or `myvacuumagent.py`) as a template implement an agent with an internal state that helps the agent clean more effectively/efficiently than the reactive/random agents. Note that an example of an agent with a simple state representation is provided (very limited - it only moves forward and sucks dirt) that you may modify in any way you wish. The agent must be able to suck up all dirt in a rectangular world of unknown dimensions **without obstacles** and shut down on the **home** position (home position can be sensed through one of the percepts). At the end the agent should also have a fully updated world model (i.e. world variable in `MyAgentState` class in Java or `self.world` in Python).  
**Note:** The agent starts from a random location which is achieved by executing `moveToRandomStartPosition()` method in the `MyvacuumAgent` file (`move_to_random_start_position` in Python). You should not change this part of the code.  
In Java you can use the already mentioned `compile` script to compile your changed agent (or use Run if you are using the Eclipse editor, files are compiled automatically when changed). In Python simply edit `myvacuumagent.py` file and run the lab again. Evaluate the performance of the agent. In what ways is it better or worse than the reactive agent?
2. In this task you will extend and improve your agent to be able to solve the problem **with obstacles** (0.1 obstacle density, 15x15 world). This can be achieved in a variety of ways with different levels of complexity (e.g. rule-based, search). Returning to the **home** position is not strictly required, but encouraged. It is sufficient that the agent performs a predefined minimum number of steps (set `iterationCounter` - Java, `self.iteration_counter` - Python variable to e.g. `[width x height x 2]`) before shutting down. Prepare to discuss during the lab demonstration in what way you have improved the agent to handle obstacles.  
**Note:** If you are already familiar with search algorithms, feel free to implement an agent based on search which shuts down at the home position after all the dirt is cleaned. In that case the `iterationCounter` should be set to a large enough value so that the agent is able to finish the task.
3. Write explanations of how your agent works (in both Task 1 and 2), what is the main decision-"loop", some notes about what data you store in your agent, and why you have chosen your solutions, if there were alternatives etc. This should add up to a report of approx. 1-2 A4s. Commented or easily readable code is also required.

## How to report your results

Demonstrate the lab to the lab assistant and hand in the code as well as the report from Task 3 according to the **submission instructions**.

Page responsible: **Fredrik Heintz**

Last updated: 2023-08-31